

Lecture 2

Divide-and-conquer, MergeSort, and Big-O notation

Today

- Things we want to know about algorithms:
 - Does it work?
 - Is it efficient?
- We'll start to see how to answer these by looking at some examples of sorting algorithms.
 - InsertionSort
 - MergeSort



SortingHatSort not discussed

The plan

- Part I: Sorting Algorithms


- InsertionSort: does it work and is it fast?
- MergeSort: does it work and is it fast?

- Skills:

- Analyzing correctness of iterative and recursive algorithms.
- Analyzing running time of recursive algorithms (part 1...more next time!)



- Part II: How do we measure the runtime of an algorithm?

- Worst-case analysis
- Asymptotic Analysis



Sorting

- Important primitive
- **For today**, we'll pretend all elements are distinct.



Benchmark: insertion sort

- Say we want to sort: $A = (6, 5, 3, 1, 8, 7, 2, 4)$
- “Algorithm”: Insert items one at a time.

Student sorting experiment (pumpkins!)

Insertion Sort Algorithm:

```
InsertionSort(A):  
  for i in [1:n]  
    current ← A[i]  
    j ← i-1  
    while j >= 0 and A[j] > current:  
      A[j+1] ← A[j]  
      j ← j-1  
    A[j+1] ← current
```

Insertion Sort Algorithm:


```
InsertionSort(A):  
  for i in [1:n]  
    current ← A[i]  
    j ← i-1  
    while j >= 0 and A[j] > current:  
      A[j+1] ← A[j]  
      j ← j-1  
    A[j+1] ← current
```

6 5 3 1 8 7 2 4

Insertion Sort

1. Does it work?

2. Is it fast?



Insertion Sort: running time

```
InsertionSort(A):  
  for i in [1:n]  
    current ← A[i]  
    j ← i-1  
    while j ≥ 0 and A[j] > current:  
      A[j+1] ← A[j]  
      j ← j-1  
    A[j+1] ← current
```


n iterations of
the outer loop

In the worst case,
about n iterations
of this inner loop

Running time scales like n^2

Insertion Sort

1. Does it work?
2. Is it fast?



- Okay, so it's pretty obvious that it works.



- **HOWEVER!** In the future it won't be so obvious, so let's take some time now to see how we would prove this rigorously.

Why does this work?

- Say you have a sorted list,

3	4	6	8
---	---	---	---

, and another element

5

.

- Insert

5

 right after the largest thing that's still smaller than

5

. (Aka, right after

4

).

- Then you get a sorted list:

3	4	5	6	8
---	---	---	---	---

So just use this logic at every step.



The first element, [6], makes up a sorted list.

So correctly inserting 4 into the list [6] means that [4,6] becomes a sorted list.




The first two elements, [4,6], make up a sorted list.

So correctly inserting 3 into the list [4,6] means that [3,4,6] becomes a sorted list.



The first three elements, [3,4,6], make up a sorted list.

So correctly inserting 8 into the list [3,4,6] means that [3,4,6,8] becomes a sorted list.



The first four elements, [3,4,6,8], make up a sorted list.

So correctly inserting 5 into the list [3,4,6,8] means that [3,4,5,6,8] becomes a sorted list.



YAY WE ARE DONE!

Recall: proof by induction

- Maintain a loop invariant.
- Proceed by induction.

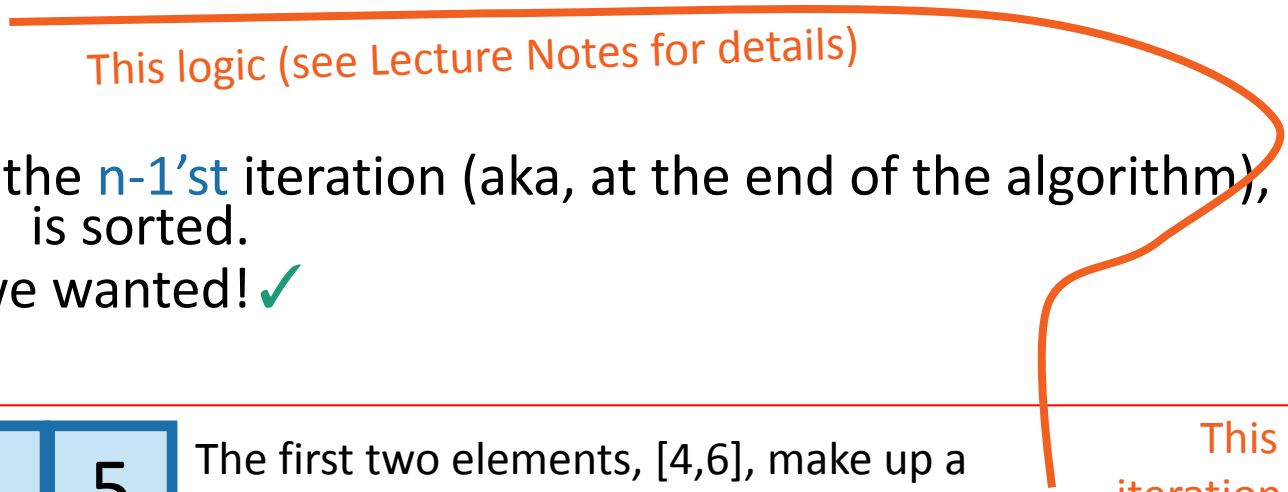
A loop invariant is something that should be true at every iteration.

- **Four steps in the proof by induction:**

- **Inductive Hypothesis:** The loop invariant holds after the i^{th} iteration.
- **Base case:** the loop invariant holds before the 1^{st} iteration.
- **Inductive step:** If the loop invariant holds after the i^{th} iteration, then it holds after the $(i+1)^{\text{st}}$ iteration
- **Conclusion:** If the loop invariant holds after the last iteration, then we win.

Formally: induction

A “loop invariant” is something that we maintain at every iteration of the algorithm.

- Loop invariant(i): $A[: i+1]$ is sorted.
- Inductive Hypothesis:
 - The loop invariant(i) holds at the end of the i^{th} iteration (of the outer loop).
- Base case ($i=0$):
 - Before the algorithm starts, $A[: 1]$ is sorted. ✓
- Inductive step: 

This logic (see Lecture Notes for details)
- Conclusion:
 - At the end of the $n-1^{\text{st}}$ iteration (aka, at the end of the algorithm), $A[: n] = A$ is sorted.
 - That's what we wanted! ✓

4	6	3	8	5
---	---	---	---	---

The first two elements, [4,6], make up a sorted list.

3	4	6	8	5
---	---	---	---	---

So correctly inserting 3 into the list [4,6] means that [3,4,6] becomes a sorted list.

This was iteration $i=2$.

Aside: proofs by induction

- We're gonna see/do/skip over a lot of them.
- I'm assuming you're comfortable with them from CS..
- If that went by too fast and was confusing:
 - Slides [there's a hidden one with more info]
 - Lecture notes
 - Book
 - Office Hours

Make sure you really understand the argument on the previous slide!



Siggi the Studious Stork

To summarize


InsertionSort is an algorithm that correctly sorts an arbitrary n -element array in time that scales like n^2 .

Can we do better?

The plan

- Part I: Sorting Algorithms

- InsertionSort: does it work and is it fast?
- MergeSort: does it work and is it fast?



- Skills:

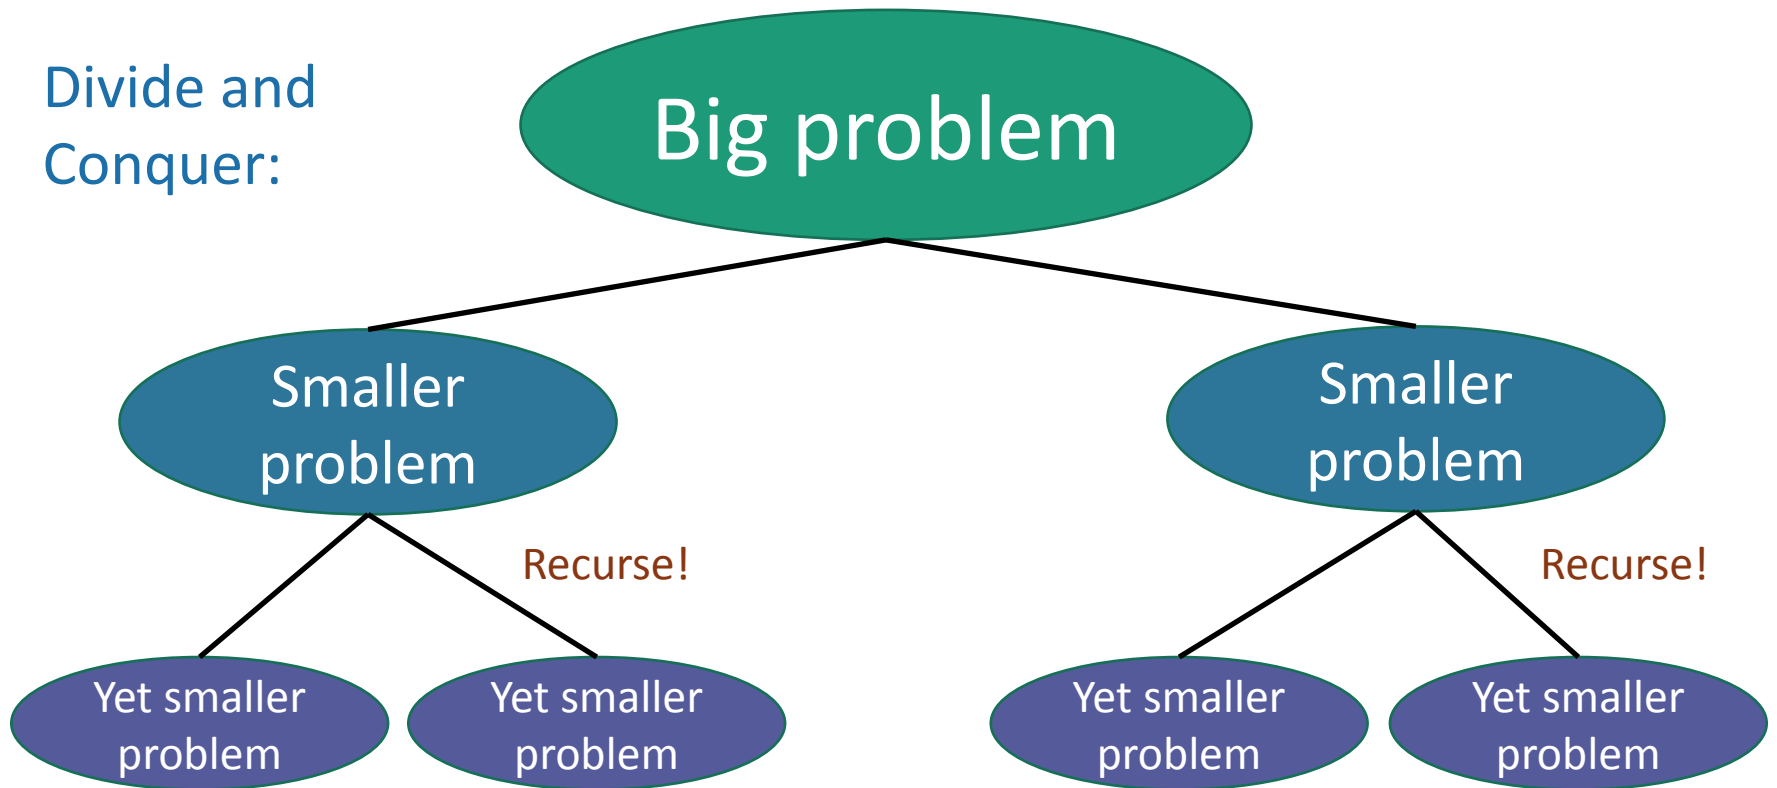
- Analyzing correctness of iterative and recursive algorithms.
- Analyzing running time of recursive algorithms (part A)

- Part II: How do we measure the runtime of an algorithm?




- Worst-case analysis
- Asymptotic Analysis

Can we do better?

- MergeSort: a divide-and-conquer approach
- Recall from last time:




MergeSort




Recursive magic!

Recursive magic!



MERGE!



How would you do this in-place?



Ollie the over-achieving Ostrich

MergeSort Pseudocode

MERGESORT(A):

- $n \leftarrow \text{length}(A)$
- **if** $n \leq 1$:
 - **return** A

If A has length 1,
It is already sorted!
- $L \leftarrow \text{MERGESORT}(A[0 : n/2])$


Sort the left half
- $R \leftarrow \text{MERGESORT}(A[n/2 : n])$

Sort the right half
- **return** **MERGE**(L,R)

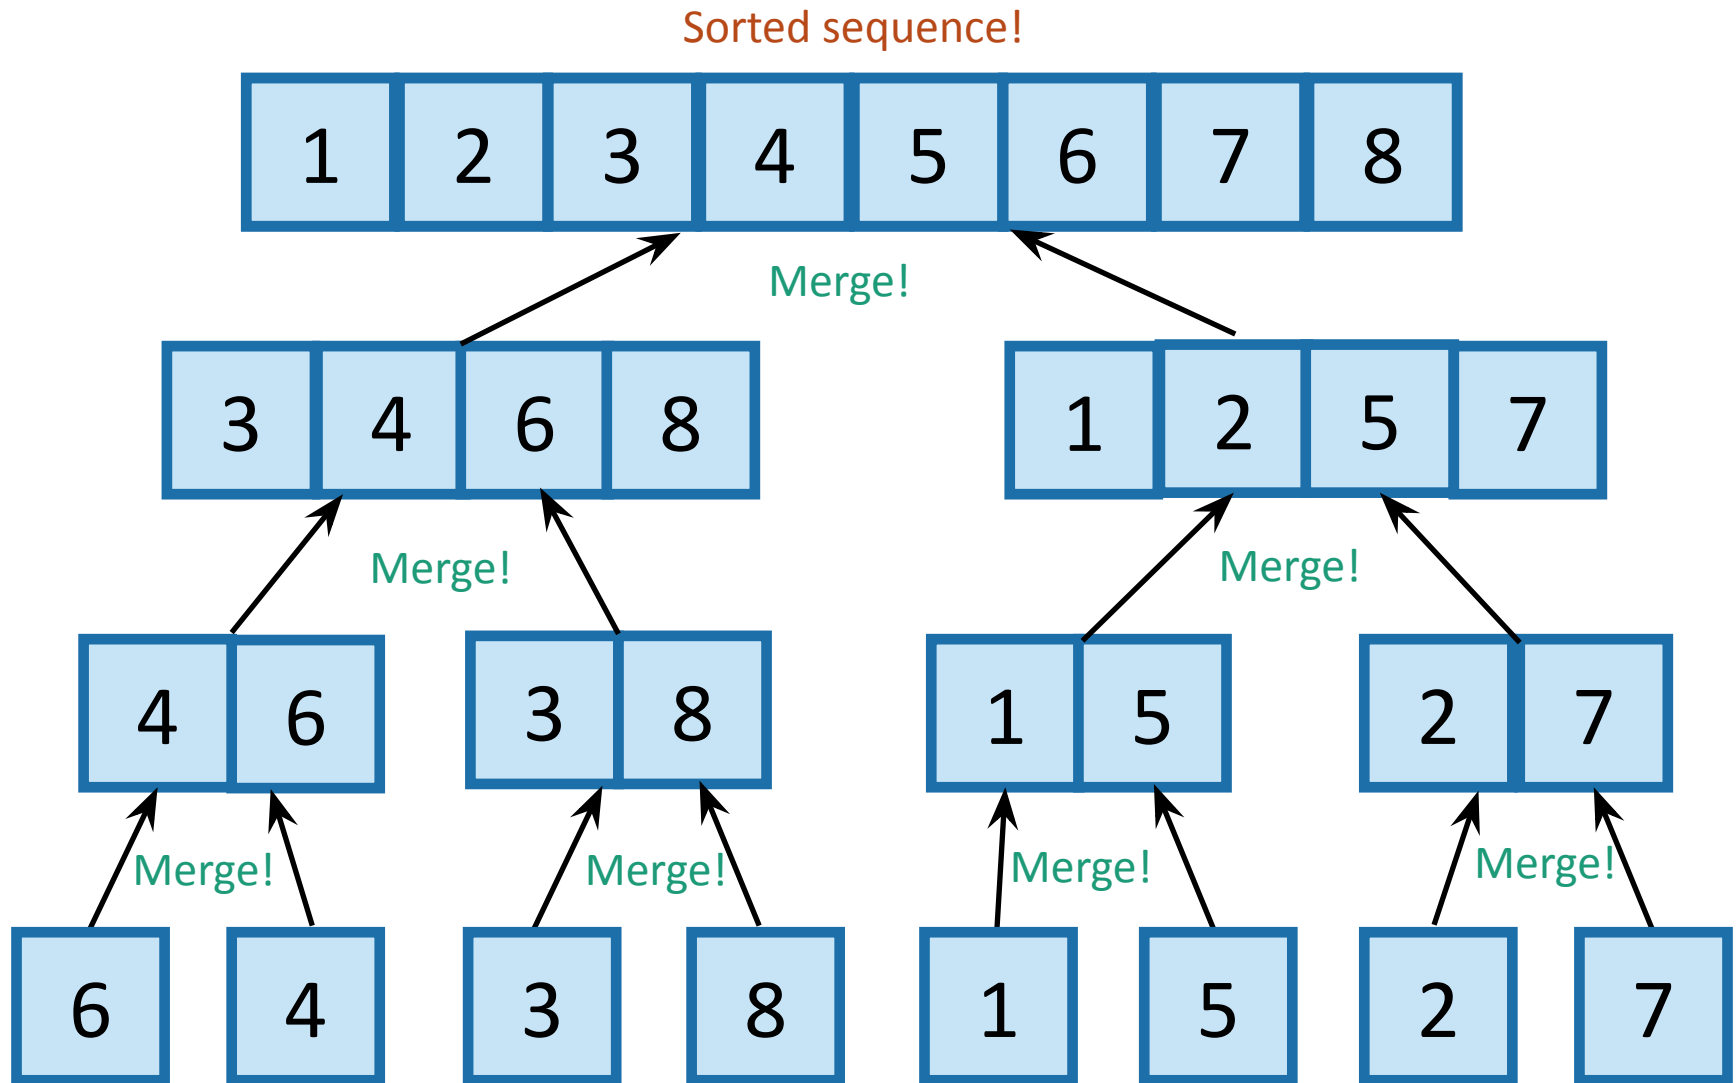
Merge the two halves

What actually happens?

First, recursively break up the array all the way down to the base cases



Then, merge them all back up!



A bunch of sorted lists of length 1 (in the order of the original sequence).

Two questions

1. Does this work?
2. Is it fast?

It works


Let's assume $n = 2^t$

Again we'll use induction.
This time with an invariant
that will remain true after
every recursive call.

- Inductive hypothesis:

“In every recursive call,
MERGESORT returns a sorted array.”

- Base case ($n=1$): a 1-element array is always sorted.
- Inductive step: Suppose that L and R are sorted. Then **MERGE**(L,R) is sorted.
- Conclusion: “In the top recursive call, **MERGESORT** returns a sorted array.”




Fill in the inductive step! (Either do it yourself or read it in CLRS!)

```
•  $n \leftarrow \text{length}(A)$ 
• if  $n \leq 1$ :
  • return A
•  $L \leftarrow \text{MERGESORT}(A[1 : n/2])$ 
•  $R \leftarrow \text{MERGESORT}(A[n/2+1 : n])$ 
• return MERGE(L,R)
```


Two questions

1. Does this work?
2. Is it fast?



Think-Pair-Share:

(2 min: try to think- how fast is MergeSort?

2 min: what does the person next to you think? why?)

MergeSort Pseudocode

MERGESORT(A):

- $n \leftarrow \text{length}(A)$
- **if** $n \leq 1$:
 - **return** A

If A has length 1,
It is already sorted!
- $L \leftarrow \text{MERGESORT}(A[0 : n/2])$

Sort the left half
- $R \leftarrow \text{MERGESORT}(A[n/2 : n])$

Sort the right half
- **return** **MERGE**(L,R)

Merge the two halves

It's fast

Let's keep assuming $n = 2^t$

CLAIM:

MERGESORT requires at most $11n (\log(n) + 1)$ operations to sort n numbers.

What exactly is an “operation” here?
We're leaving that vague on purpose.
Also I made up the number 11.



How does this compare to
InsertionSort?

Scaling like n^2 vs scaling like $n \log(n)$?

Quick log refresher

All logarithms in this course are base 2

- $\log(n)$: how many times do you need to divide n by 2 in order to get down to 1?

32

16

8

4

2

1

$$\log(32) = 5$$

64

32

16

8

4

2

1

$$\log(64) = 6$$

$$\log(128) = 7$$

$$\log(256) = 8$$

$$\log(512) = 9$$

.

.

.

$$\log(\text{number of particles in the universe}) < 280$$

Moral: $\log(n)$ grows very slowly with n .

It's fast!


CLAIM:

MERGESORT requires at most $11n (\log(n) + 1)$ operations to sort n numbers.

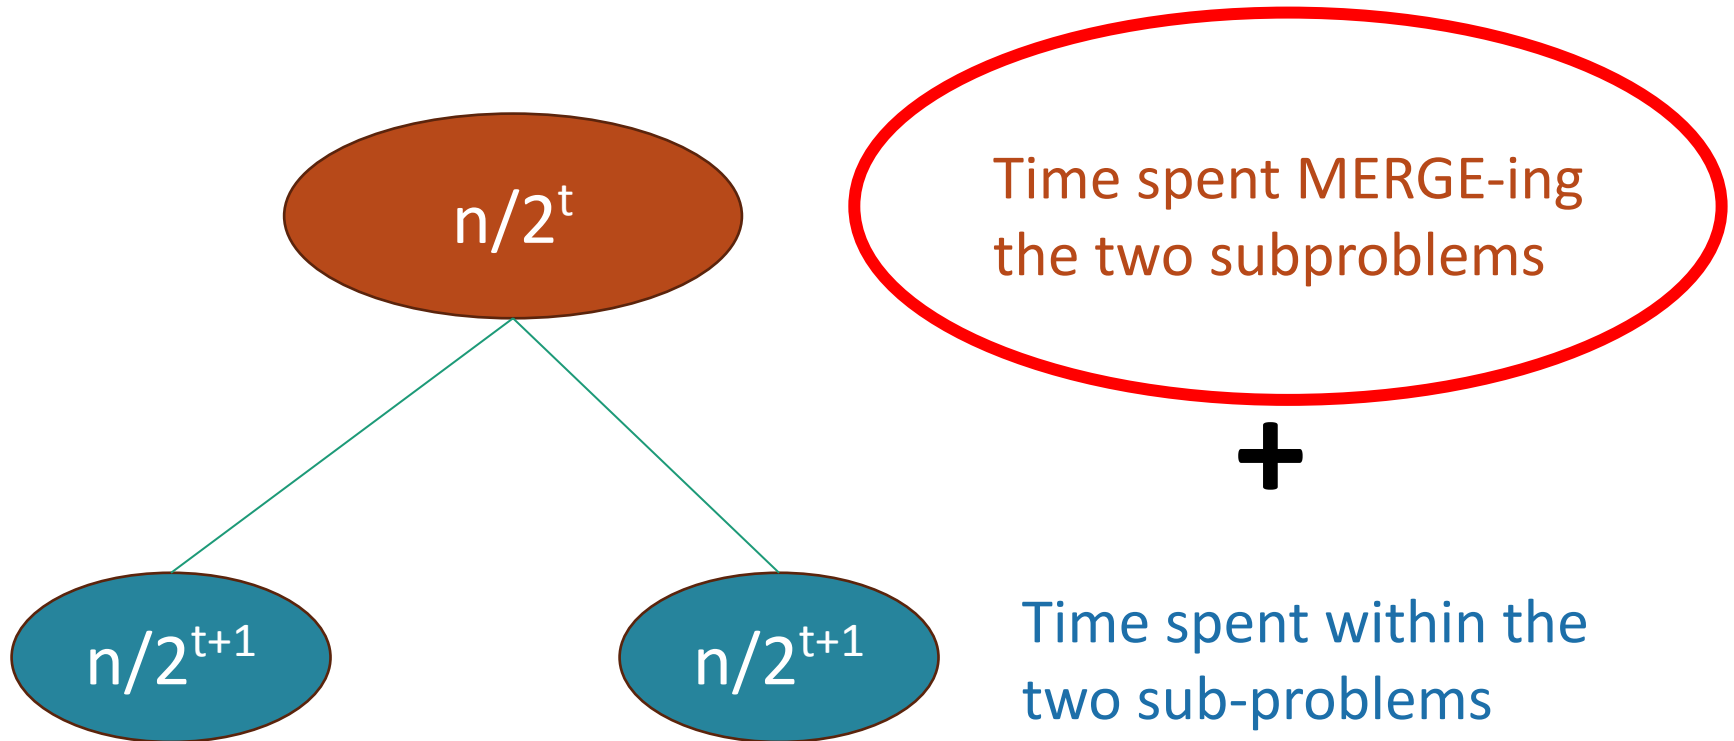
Much faster than InsertionSort for large n !

Let's prove the claim

- Later we'll see **more principled** ways of analyzing divide-and-conquer algs.
- **But for today let's just wing it.**

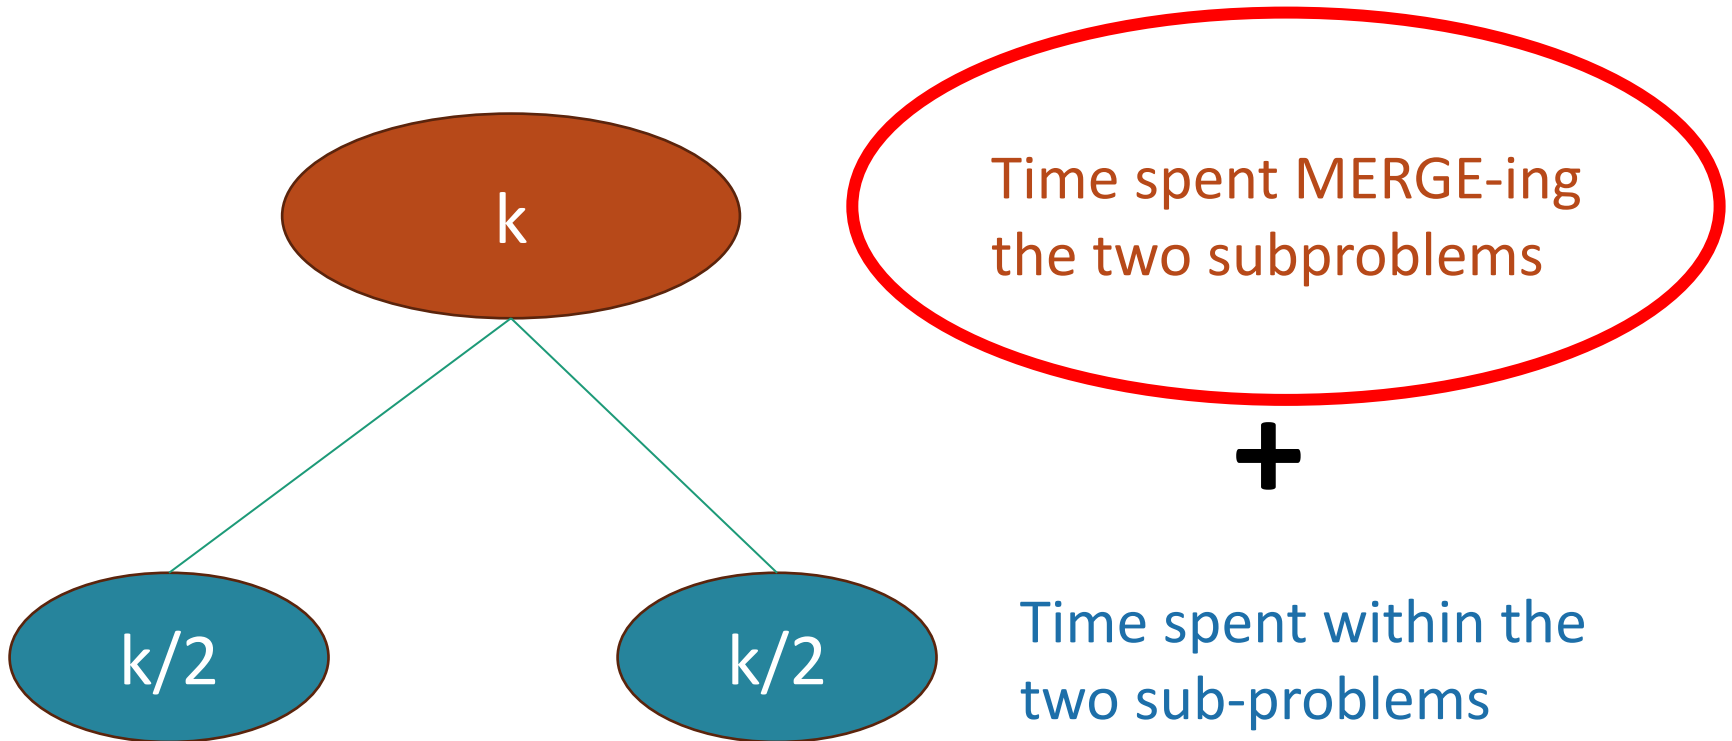


How much work in this sub-problem?






How much work in this sub-problem?


Let $k=n/2^t \dots$




How long does it take to MERGE?



MERGE!



How long does it take to MERGE?



- Time to initialize an array of size k
- Plus the time to initialize three counters
- Plus the time to increment two of those counters $k/2$ times each
- Plus the time to compare two values at least k times
- Plus the time to copy k values from the existing array to the big array.
- Plus...

Let's say no more than **11k** operations.

There's some justification for this number "11" in the lecture notes, but it's really pretty arbitrary.

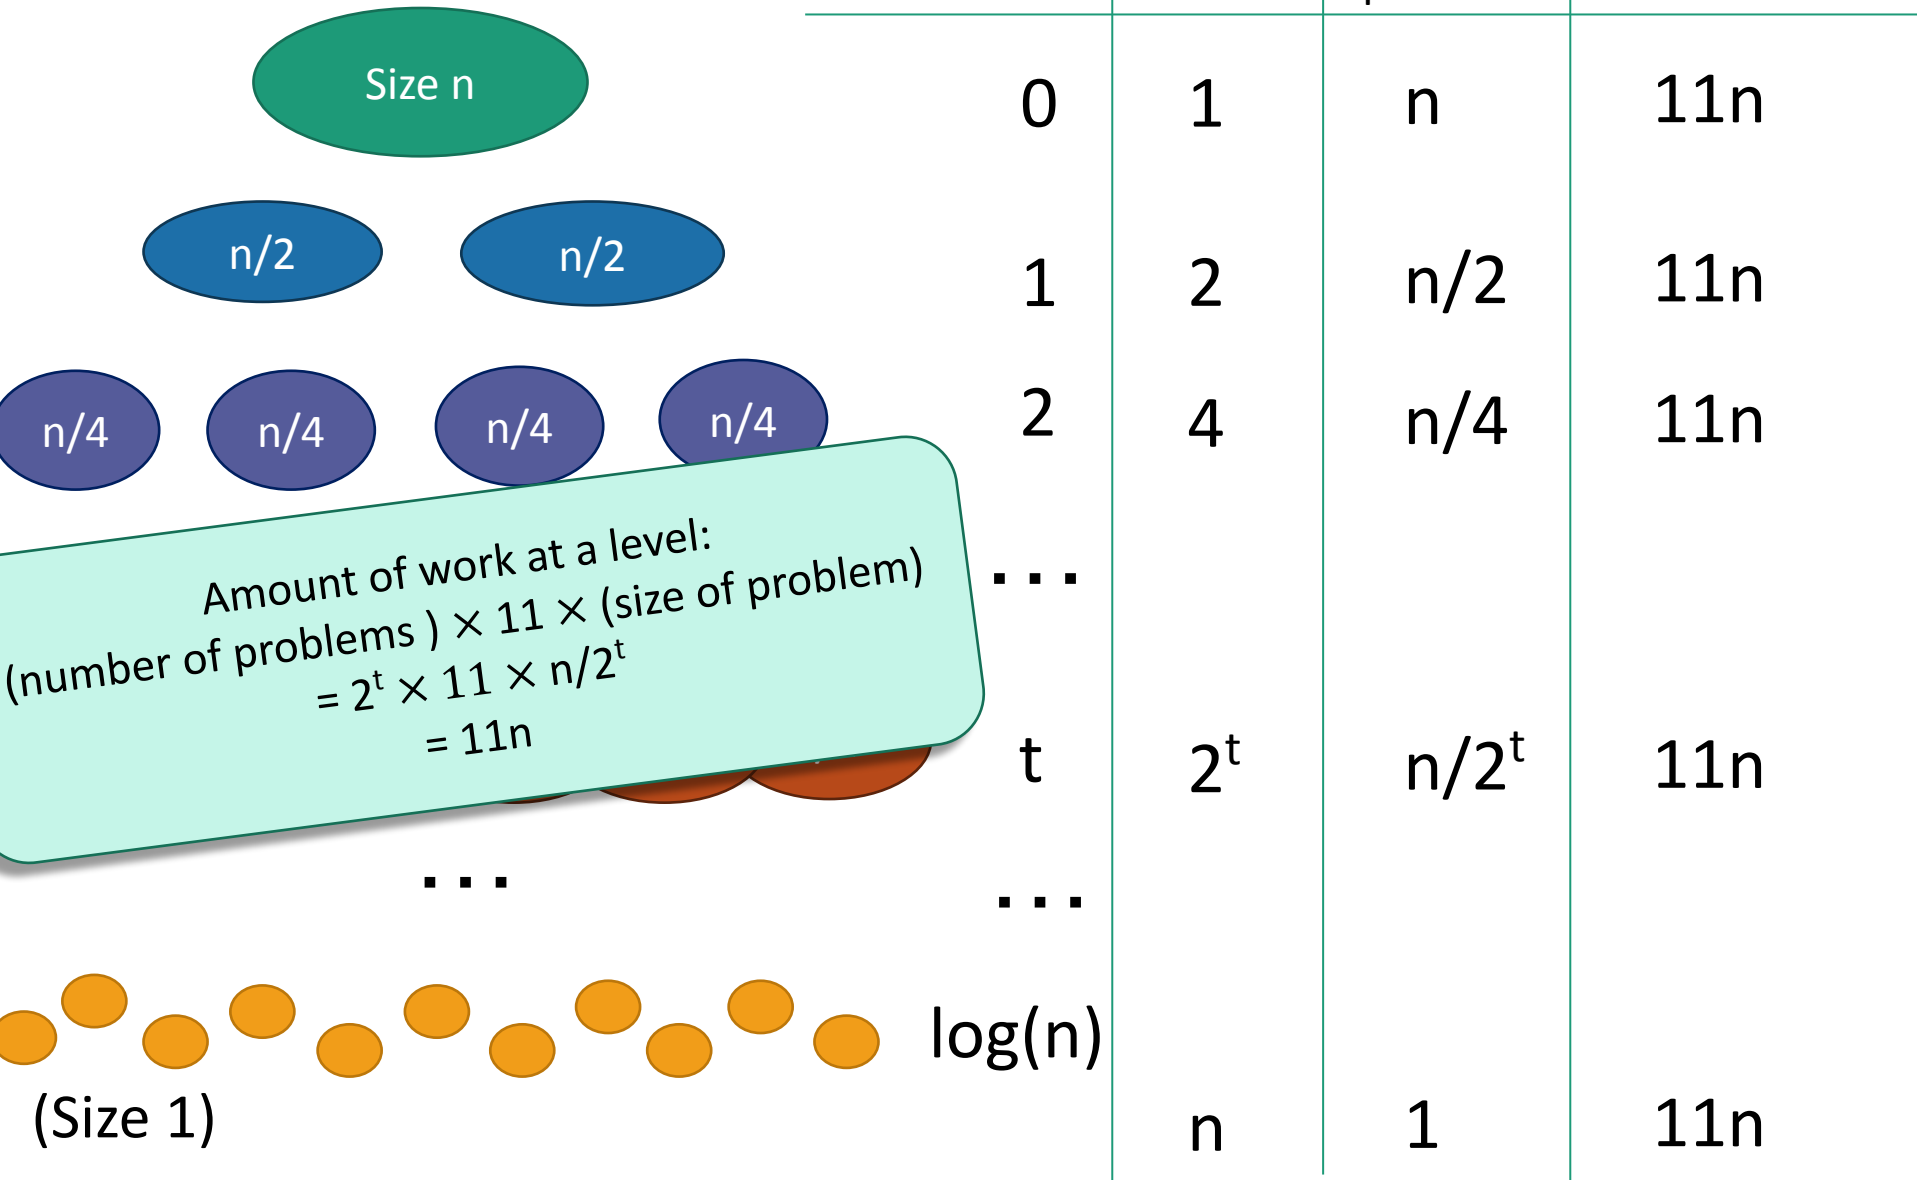


Plucky the
Pedantic Penguin



Lucky the
lackadaisical lemur

Recursion tree




Total runtime...

- $11n$ steps per level, at every level
- $\log(n) + 1$ levels
- $11n (\log(n) + 1)$ steps total

That was the claim!

A few reasons to be grumpy

- Sorting




should take zero steps...

- What's with this 11k bound?
 - You made that number "11" up.
 - Different operations don't take the same amount of time.



How we will deal with **grumpiness**

- Take a deep breath...
- Worst case analysis
- Asymptotic notation



The plan

- Part I: Sorting Algorithms

- InsertionSort: does it work and is it fast?
- MergeSort: does it work and is it fast?
- Skills:
 - Analyzing correctness of iterative and recursive algorithms.
 - Analyzing running time of recursive algorithms (part A)



- Part II: How do we measure the runtime of an algorithm?


- Worst-case analysis
- Asymptotic Analysis

Worst-case analysis

Sorting a sorted list
should be fast!!

1	2	3	4	5	6	7	8
---	---	---	---	---	---	---	---

- In this class, we will focus on **worst-case analysis**



Algorithm
designer

Here is my algorithm!

```
Algorithm:  
  Do the thing  
  Do the stuff  
  Return the answer
```

Here is an input!



- **Pros:** very strong guarantee
- **Cons:** very strong guarantee

Why worst-case analysis?

The real reasons:

1. We don't really know anything much better
 - Very popular these days: “average case analysis”
 - Downside: we typically don't know what an average input looks like.
2. Best-case + worst-case \neq average-case

$$\begin{array}{lcl} \text{Best-case} & + & \text{worst case} = \text{worst-case} \\ O(1) & + & O(n \log n) = O(n \log n) \end{array}$$



Big-O notation

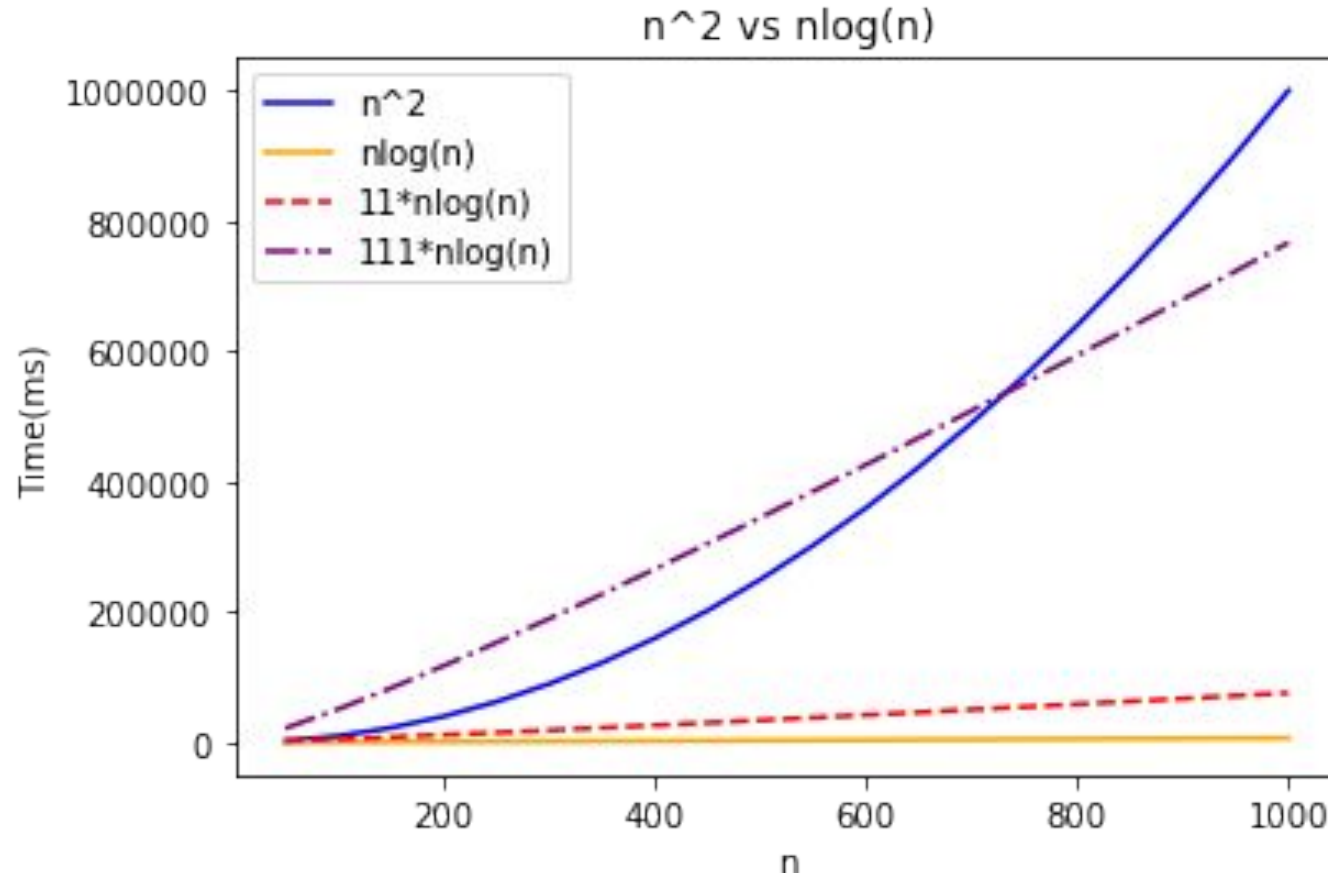
How long does an operation take? Why are we being so sloppy about that “11”?



- What do we mean when we measure runtime?
 - We probably care about wall time: how long does it take to solve the problem, in seconds or minutes or hours?
- This is heavily dependent on the programming language, architecture, etc.
- These things are very important, but are **not the point of this class.**
- We want a way to talk about the running time of an algorithm, **independent of these considerations.**

Main idea:

Focus on how the runtime **scales** with n (the input size).



Asymptotic Analysis

How does the running time scale as n gets large?

One algorithm is “faster” than another if its runtime scales better with the size of the input.

Pros:

- Abstracts away from hardware- and language-specific issues.
- Makes algorithm analysis much more tractable.

Cons:

- Only makes sense if n is large (compared to the constant factors).

$2^{1000000000000000} n$
is “better” than n^2 ?!?!

pronounced “big-oh of ...” or sometimes “oh of ...”

$O(\dots)$ means an upper bound

- Let $T(n)$, $g(n)$ be positive functions of positive integers.
 - Think of $T(n)$ as being a runtime: positive and increasing in n .
- We say “ $T(n)$ is $O(g(n))$ ” if $g(n)$ grows at least as fast as $T(n)$ as n gets large.
- Formally,

$$\begin{aligned} T(n) &= O(g(n)) \\ &\iff \\ \exists c, n_0 > 0 \text{ s.t. } \forall n \geq n_0, \\ 0 &\leq T(n) \leq c \cdot g(n) \end{aligned}$$

Example


$$2n^2 + 10 = O(n^2)$$

$$T(n) = O(g(n))$$

$$\Leftrightarrow$$

$$\exists c, n_0 > 0 \text{ s.t. } \forall n \geq n_0,$$

$$0 \leq T(n) \leq c \cdot g(n)$$



Example


$$2n^2 + 10 = O(n^2)$$

$$T(n) = O(g(n))$$

$$\Leftrightarrow$$

$$\exists c, n_0 > 0 \text{ s.t. } \forall n \geq n_0,$$

$$0 \leq T(n) \leq c \cdot g(n)$$



Formally:

- Choose $c = 3$
- Choose $n_0 = 4$
- Then:

$$\forall n \geq 4,$$

$$0 \leq 2n^2 + 10 \leq 3 \cdot n^2$$


same Example
 $2n^2 + 10 = O(n^2)$

$$T(n) = O(g(n))$$

$$\Leftrightarrow$$

$$\exists c, n_0 > 0 \text{ s.t. } \forall n \geq n_0,$$

$$0 \leq T(n) \leq c \cdot g(n)$$



Formally:

- Choose $c = 7$
- Choose $n_0 = 2$
- Then:


$$\forall n \geq 2,$$

$$0 \leq 2n^2 + 10 \leq 7 \cdot n^2$$

There isn't a
unique "correct"
choice of c and n_0

Another example:
 $n = O(n^2)$

$$T(n) = O(g(n))$$
$$\Leftrightarrow \exists c, n_0 > 0 \text{ s.t. } \forall n \geq n_0,$$
$$0 \leq T(n) \leq c \cdot g(n)$$



- Choose $c = 1$
- Choose $n_0 = 1$
- Then

$$\forall n \geq 1,$$
$$0 \leq n \leq n^2$$

$\Omega(\dots)$ means a lower bound

- We say “ $T(n)$ is $\Omega(g(n))$ ” if $T(n)$ grows at least as fast as $g(n)$, as n gets large.

- Formally,

$$T(n) = \Omega(g(n))$$

$$\Leftrightarrow$$

$$\exists c, n_0 > 0 \text{ s.t. } \forall n \geq n_0,$$

$$0 \leq c \cdot g(n) \leq T(n)$$


Switched these!!

Example


$n \log_2(n) = \Omega(3n)$

$$T(n) = \Omega(g(n))$$

$$\Leftrightarrow$$

$$\exists c, n_0 > 0 \text{ s.t. } \forall n \geq n_0,$$

$$0 \leq c \cdot g(n) \leq T(n)$$



- Choose $c = 1/3$
- Choose $n_0 = 2$
- Then

$$\forall n \geq 2,$$

$$0 \leq \frac{3n}{3} \leq n \log_2(n)$$

$\Theta(\dots)$ means both!

- We say “ $T(n)$ is $\Theta(g(n))$ ” if:

$$T(n) = O(g(n))$$

-AND-

$$T(n) = \Omega(g(n))$$

Some more examples

- All degree k polynomials are $O(n^k)$
- For any $k \geq 1$, n^k is **not** $O(n^{k-1})$

(On the board if we have
time... if not see the lecture
notes!)

Take-away from examples


- To prove $T(n) = O(g(n))$, you have to come up with c and n_0 so that the definition is satisfied.
- To prove $T(n)$ is **NOT** $O(g(n))$, one way is **proof by contradiction**:
 - Suppose (to get a contradiction) that someone gives you a c and an n_0 so that the definition *is* satisfied.
 - Show that this someone must be lying to you by deriving a contradiction.

Some brainteasers

- Are there functions f, g so that **NEITHER** $f = O(g)$ nor $f = \Omega(g)$?
- Are there **non-decreasing** functions f, g so that the above is true?
- Define the n 'th fibonacci number by $F(0) = 1, F(1) = 1, F(n) = F(n-1) + F(n-2)$ for $n > 2$.
 - 1, 1, 2, 3, 5, 8, 13, 21, 34, 55, ...

True or false:

- $F(n) = O(2^n)$
- $F(n) = \Omega(2^n)$



Ollie the Over-achieving Ostrich

What have we learned?

Asymptotic Notation

- This makes both Plucky and Lucky happy.
 - **Plucky the Pedantic Penguin** is happy because there is a precise definition.
 - **Lucky the Lackadaisical Lemur** is happy because we don't have to pay close attention to all those pesky constant factors like "11".
- But we should be careful not to abuse it.
- In this course, (almost) every algorithm we see will be actually practical, without needing to take $n \geq n_0 = 2^{100000000}$.

This is my
happy face!



The plan

- Part I: Sorting Algorithms

- InsertionSort: does it work and is it fast?
- MergeSort: does it work and is it fast?

- Skills:

- Analyzing correctness of iterative and recursive algorithms.
- Analyzing running time of recursive algorithms (part A)

- Part II: How do we measure the runtime of an algorithm?

- Worst-case analysis
- Asymptotic Analysis

Wrap-Up 

Recap

- InsertionSort runs in time $O(n^2)$
- MergeSort is a divide-and-conquer algorithm that runs in time $O(n \log(n))$
- How do we show an algorithm is correct?
 - Today, we did it by induction
- How do we measure the runtime of an algorithm?
 - Worst-case analysis
 - Asymptotic analysis

Next time

- A more systematic approach to analyzing the runtime of recursive algorithms.